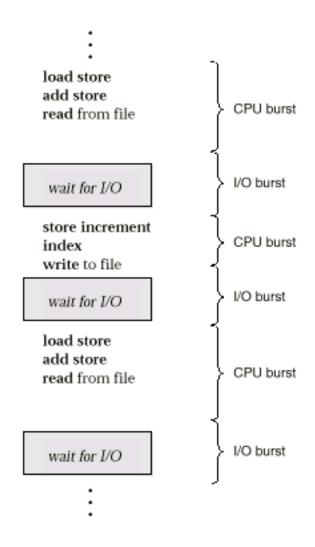
Module 6: CPU Scheduling

- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms
- Multiple-Processor Scheduling
- Real-Time Scheduling
- Algorithm Evaluation

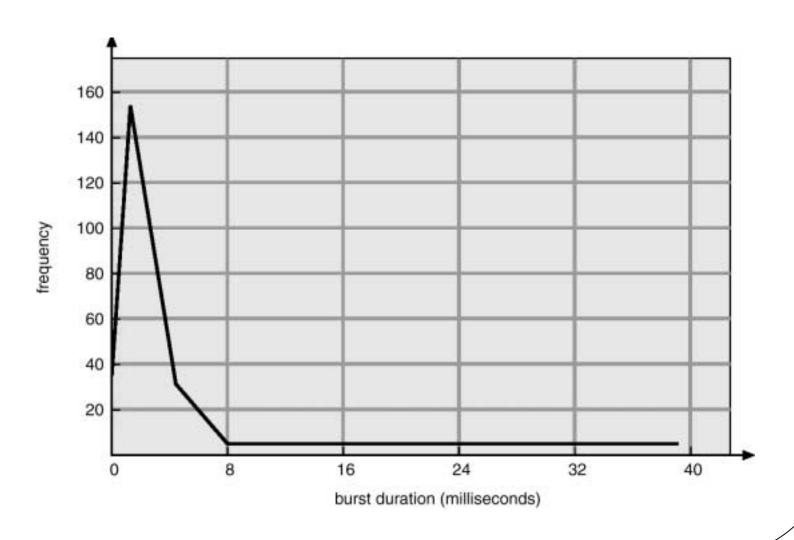
Basic Concepts

- Maximum CPU utilization obtained with multiprogramming
- CPU-I/O Burst Cycle Process execution consists of a cycle of CPU execution and I/O wait.
- CPU burst distribution

Alternating Sequence of CPU And I/O Bursts



Histogram of CPU-burst Times



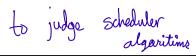
CPU Scheduler

- Selects from among the processes in memory that are ready to execute, and allocates the CPU to one of them.
- CPU scheduling decisions may take place when a process:
 - 1. Switches from running to waiting state.
 - 2. Switches from running to ready state.
 - 3. Switches from waiting to ready.
 - 4. Terminates.
- Scheduling under 1 and 4 is nonpreemptive.
- All other scheduling is preemptive.

Dispatcher

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - switching context
 - switching to user mode
 - jumping to the proper location in the user program to restart that program
- Dispatch latency time it takes for the dispatcher to stop one process and start another running.

Scheduling Criteria to judge scheduler algaritim



- CPU utilization keep the CPU as busy as possible
- Throughput # of processes that complete their execution per time unit
- Turnaround time amount of time to execute a particular process
- Waiting time amount of time a process has been wiating in the ready queue for the CPU (not just first)
- Response time amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

Optimization Criteria

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time

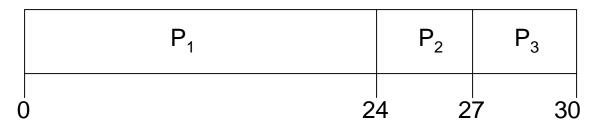
First-Come, First-Served (FCFS) Scheduling

• Example:

 $\frac{\text{Process}}{P_1} \qquad \frac{\text{estimated}}{\text{Burst Time}}$ $\frac{P_2}{P_2} \qquad 3$

3

• Suppose that the processes arrive in the order: P_1 , P_2 , P_3 The Gantt Chart for the schedule is:



- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: (0 + 24 + 27)/3 = 17

FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order

$$P_2$$
, P_3 , P_1 .

• The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: (6 + 0 + 3)/3 = 3
- Much better than previous case.
- Convoy effect short process behind long process

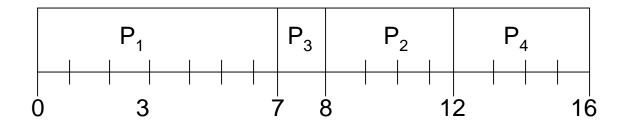
Shortest-Job-First (SJR) Scheduling

- Associate with each process the length of its next CPU burst.
 Use these lengths to schedule the process with the shortest time.
- Two schemes:
 - nonpreemptive once CPU given to the process it cannot be preempted until completes its CPU burst.
 - Preemptive if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This scheme is know as the Shortest-Remaining-Time-First (SRTF).
- SJF is optimal gives minimum average waiting time for a given set of processes.

Example of Non-Preemptive SJF

<u>Process</u>	Arrival Time	Burst Time
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

• SJF (non-preemptive)




• Average waiting time = (0 + 6 + 3 + 7)/4 - 4

Example of Preemptive SJF

<u>Process</u>	Arrival Time	Burst Time
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

• SJF (preemptive)



• Average waiting time = (9 + 1 + 0 + 2)/4 - 3

Determining Length of Next CPU Burst

- Can only estimate the length.
- Can be done by using the length of previous CPU bursts, using exponential averaging.
 - 1. $t_n = \text{actual lenght of } n^{th} \text{CPU burst}$
 - 2. $\tau_{n+1} = \text{predicted value for the next CPU burst}$ 3. $\alpha, 0 \le \alpha \le 1$ weight for adval vs. predicted

 - 4. Define:

$$au_{n=1} = \alpha t_n + (1-\alpha)\tau_n.$$

Examples of Exponential Averaging

- \bullet $\alpha = 0$
 - $\tau_{n+1} = \tau_n$
 - Recent history does not count.
- \bullet $\alpha = 1$
 - $\tau_{n+1} = t_n$
 - Only the actual last CPU burst counts.
- If we expand the formula, we get:

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \alpha t_n - 1 + \dots + (1 - \alpha)^j \alpha t_n - 1 + \dots + (1 - \alpha)^{n+1} t_n \tau_0$$

• Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor.

Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer = highest priority).
 - Preemptive
 - nonpreemptive
- SJF is a priority scheduling where priority is the predicted next CPU burst time.
- Problem = Starvation low priority processes may never execute.
- Solution = Aging as time progresses increase the priority of the process.

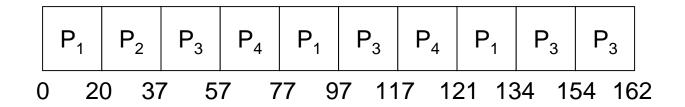
Round Robin (RR)

- Each process gets a small unit of CPU time (time quantum), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q, then each process gets 1/n of the CPU time in chunks of at most q time units at once. No process waits more than (n-1)q time units.
- Performance
 - q large \Rightarrow FIFO
 - q small ⇒ q must be large with respect to context switch, otherwise overhead is too high.

Example: RR with Time Quantum = 20

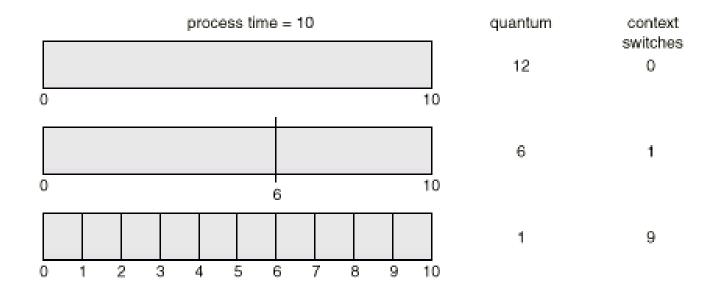
<u>Process</u>	Burst Time
P_1	53
P_2	17
P_3	68
P_4	24

• The Gantt chart is:

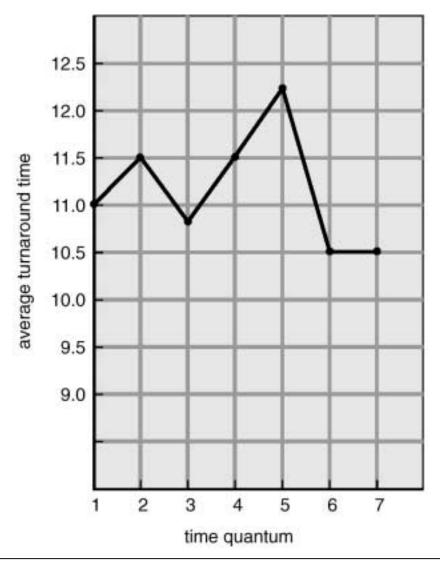


• Typically, higher average turnaround than SJF, but better response.

How a Smaller Time Quantum Increases Context Switches



Turnaround Time Varies With The Time Quantum

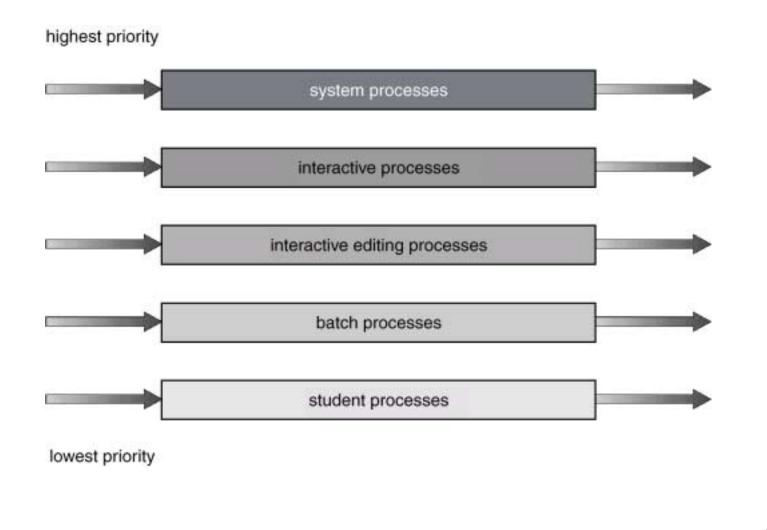


process	time
P ₁	6
P_2	3
P_3	1
P_4	7

Multilevel Queue

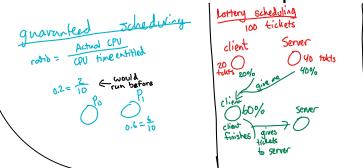
- Ready queue is partitioned into separate queues: foreground (interactive) background (batch)
- Each queue has its own scheduling algorithm, foreground – RR background – FCFS
- Scheduling must be done between the queues.
 - Fixed priority scheduling; i.e., serve all from foreground then from background. Possibility of starvation.
 - Time slice each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR
 - 20% to background in FCFS

Multilevel Queue Scheduling

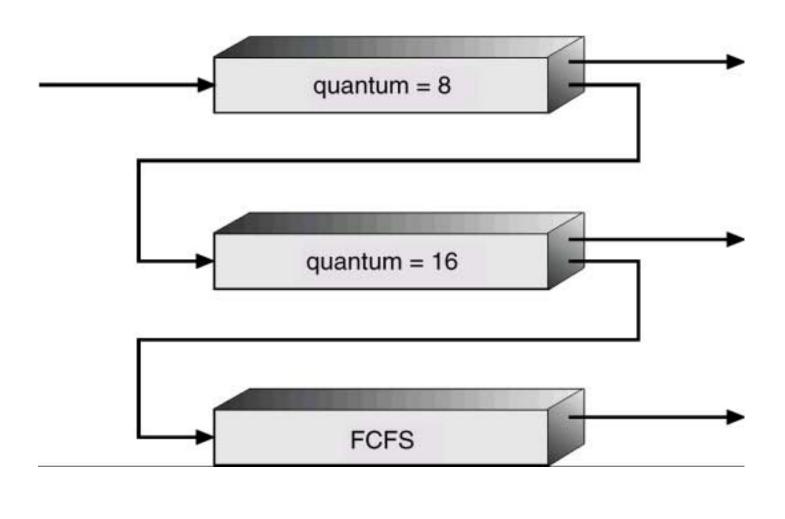


Multilevel Feedback Queue

- A process can move between the various queues; aging can be implemented this way.
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service



Multilevel Feedback Queues



Example of Multilevel Feedback Queue

• Three queues:

- Q₀ time quantum 8 milliseconds
- Q₁ time quantum 16 milliseconds
- $-Q_2 FCFS$

Scheduling

- A new job enters queue Q₀ which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue Q₁.
- At Q₁ job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q₂.

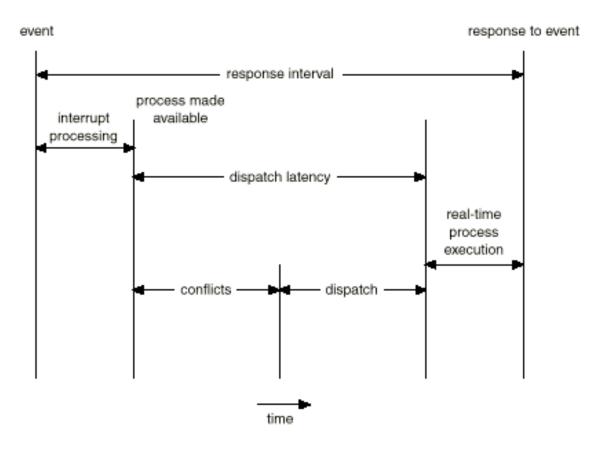
Multiple-Processor Scheduling

- CPU scheduling more complex when multiple CPUs are available.
- Homogeneous processors within a multiprocessor.
- Load sharing
- Symmetric Multiprocessing (SMP) each processor makes its own scheduling decisions.
- Asymmetric multiprocessing only one processor accesses the system data structures, alleviating the need for data sharing.

Real-Time Scheduling

- Hard real-time systems required to complete a critical task within a guaranteed amount of time.
- Soft real-time computing requires that critical processes receive priority over less fortunate ones.

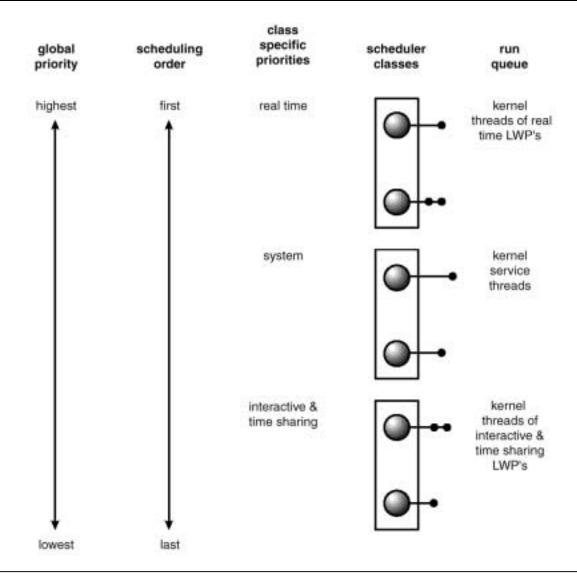
Dispatch Latency



Thread Scheduling

- Local Scheduling How the threads library decides which thread to put onto an available LWP.
- Global Scheduling How the kernel decides which kernel thread to run next.

Solaris 2 Scheduling



Java Thread Scheduling

JVM Uses a Preemptive, Priority-Based Scheduling Algorithm.

 FIFO Queue is Used if There Are Multiple Threads With the Same Priority.

Java Thread Scheduling (cont)

JVM Schedules a Thread to Run When:

- The Currently Running Thread Exits the Runnable State.
- A Higher Priority Thread Enters the Runnable State
- * Note the JVM Does Not Specify Whether Threads are Time-Sliced or Not.

Time-Slicing

 Since the JVM Doesn't Ensure Time-Slicing, the yield() Method May Be Used:

```
while (true) {
    // perform CPU-intensive task
    . . .
    Thread.yield();
}
```

This Yields Control to Another Thread of Equal Priority.

Thread Priorities

Thread Priorities:

<u>Priority</u>	<u>Comment</u>
Thread.MIN_PRIORITY	Minimum Thread Priority
Thread.MAX_PRIORITY	Maximum Thread Priority
Thread.NORM_PRIORITY	Default Thread Priority

Priorities May Be Set Using setPriority() method: setPriority(Thread.NORM_PRIORITY + 2);

Algorithm Evaluation

- Deterministic modeling takes a particular predetermined workload and defines the performance of each algorithm for that workload.
- Queuing models
- Implementation

Evaluation of CPU Schedulers by Simulation

